



SKILLPILLS

Skill Pill: Julia

Lecture 4: Distributed and parallel computing

Valentin Churavy James Schloss

Okinawa Institute of Science and Technology

valentin.churavy@oist.jp

james.schloss@oist.jp

July 13, 2017



OIST

- 1 Levels of parallelism
- 2 Instruction level parallelism
- 3 Threading
- 4 Distributed
- 5 CUDAnative

Necessary packages

- ▶ SIMD.jl
- ▶ MPI.jl
- ▶ DistributedArrays.jl
- ▶ CUDAnative.jl if your computer has a NVidia GPU

- ▶ Instruction level parallelism
- ▶ Shared-memory and threading
- ▶ Distributed
- ▶ Accelerators e.g. GPGPU

```
function padd(a, b, x, y)
    c = a + b
    z = x + y
    return c, z
end
```

Observation

The two operations are independent of each other and we could execute them in parallel.

1. Use `@code_llvm` and `@code_native` to understand what is happening
2. Establish a baseline performance with `@benchmark`
3. Start Julia with `julia -O3`
4. Compare the llvm and native code and your benchmark results
5. Note that there is next to no performance benefit in this example, but that changes once you scale up

```
function add(out, x, y)
    for i in 1:length(out)
        out[i] = x[i] + y[i]
    end
end
```

Observation

Each loop iteration is independent.

1. Learn about `@inbounds`

```
function sum(x)
    acc = 0.0
    for i in 1:length(x)
        acc += x[i]
    end
end
```

Observation

Is each loop iteration independent from each other? Yes and no. Standard addition is associative and the order of operations has no impact. Floating-point addition is non-associative and the order of operations is important. The compiler cannot vectorise this loop, without changing the semantics.

1. Learn about @simd

SIMD.jl

Instead of relying on the compiler to optimise and vectorise our code correctly we can also write explicit SIMD code.

```
using SIMD
function add(out::Vector{Float64}, x::Vector{Float64},
            y::Vector{Float64})
    # My laptop supports AVX 256bit 4xFloat64
    @assert length(x) % 4 == 0
    for i in 1:4:length(x)
        vx = vload(Vec{4, Float64}, x, i)
        vy = vload(Vec{4, Float64}, y, i)
        vo = vx + vy
        vstore(vo, out, i)
    end
end
```

Caveat

Julia threading model is based on a fork-join approach and is still considered experimental. The Julia developers are working with Intel on bringing a modern and full-fledged threading model to Julia for 1.0

Caveat

Julia threading model is based on a fork-join approach and is still considered experimental. The Julia developers are working with Intel on bringing a modern and full-fledged threading model to Julia for 1.0

What is fork-join

Fork-join describes the control flow that a group of threads undergo. Execution is forked across all threads and at a later point execution is joined together. In Julia all threads have to execute the same lambda.

Start Julia with `JULIA_NUM_THREADS=4` julia the number of threads needs to be set before starting Julia.

```
using Base.Threads

a = zeros(nthreads()*10)
@threads for i in 1:length(a)
    a[i] = threadid()
end
```

Caveats

Special care needs to be taken for access to global state, which includes IO and random numbers.

```
A = zeros(10_000)
@threads for i in 1:length(A)
    A[i] = rand() # Note rand uses a global variable
                Base.GLOBAL_RNG
end

rngs = [MersenneTwister(rand(UInt64)) for i in 1:nthreads()]
@threads for i in 1:length(A)
    A[i] = rand(rngs[threadid()])
end
```

```
acc = 0
@threads for i in 1:10_000
    acc += 1
end
```

To prevent the data race we can use Atomics to ensure read-write consistency.

```
acc = Atomic{Int64}(0)
@threads for i in 1:10_000
    atomic_add!(acc, 1)
end
```

Julia supports MPI through `MPI.jl`. If you are comfortable with MPI you can use it with very low overhead.

Starting Julia with `julia -p 4` will create one master process and four worker processes. You can use `remotecalls` to execute work on a worker and to get back the results. Based on this `DistributeArrays.jl` provides a powerful global array interface. `@everywhere` will execute a piece of code on all processes.

CUDAnative

A recent development is the ability to write raw CUDA code in Julia!
Learn more about it at <https://github.com/JuliaGPU/CUDAnative.jl>