



SKILLPILLS

Skill Pill: Julia

Lecture 4: Distributed and parallel computing

Valentin Churavy James Schloss

Okinawa Institute of Science and Technology

valentin.churavy@oist.jp

james.schloss@oist.jp

July 13, 2017



1 Levels of parallelism

2 Instruction level parallelism

3 Threading

4 Distributed

Necessary packages

- ▶ SIMD.jl
- ▶ MPI.jl
- ▶ DistributedArrays.jl
- ▶ CUDAnative.jl if your computer has a NVidia GPU

- ▶ Instruction level parallelism
- ▶ Shared-memory and threading
- ▶ Distributed
- ▶ Accelerators e.g. GPGPU

```
function padd(a, b, x, y)
    c = a + b
    z = x + y
    return c, z
end
```

Observation

The two operations are independent of each other and we could execute them in parallel.

1. Use `@code_llvm` and `@code_native` to understand what is happening
2. Establish a baseline performance with `@benchmark`
3. Start Julia with `julia -O3`
4. Compare the llvm and native code and your benchmark results
5. Note that there is next to no performance benefit in this example, but that changes once you scale up

```
function add(out, x, y)
    for i in 1:length(out)
        out[i] = x[i] + y[i]
    end
end
```

Observation

Each loop iteration is independent.

1. Learn about `@inbounds`

```
function sum(x)
    acc = 0.0
    for i in 1:length(x)
        acc += x[i]
    end
end
```

Observation

Is each loop iteration independent from each other? Yes and no. Standard addition is associative and the order of operations has no impact. Floating-point addition is non-associative and the order of operations is important. The compiler cannot vectorise this loop, without changing the semantics.

1. Learn about @simd

SIMD.jl

Instead of relying on the compiler to optimise and vectorise our code correctly we can also write explicit SIMD code.

```
using SIMD
function add(out::Vector{Float64}, x::Vector{Float64},
            y::Vector{Float64})
    # My laptop supports AVX 256bit 4xFloat64
    @assert length(x) % 4 == 0
    for i in 1:4:length(x)
        vx = vload(Vec{4, Float64}, x, i)
        vy = vload(Vec{4, Float64}, y, i)
        vo = vx + vy
        vstore(vo, out, i)
    end
end
```

Explain current fork join model and caveats

Simple example

complex example with loop splitting, random etc...

Atomics

MPI

Distributed Arrays